

# Overview of hidden Markov models

ISTA 410 / INFO 510: Bayesian Modeling and Inference

---

U. of Arizona School of Information

November 10, 2021

Last time:

- Linear dynamical systems and the Kalman filter

Today:

- Overview of hidden Markov models

HW5 (covarying parameters in a GLM) on D2L

# Temporal and dynamical models

---


# Temporal and dynamical models

The models we'll look at next are used to model sequential data, especially time series.

- Hidden Markov model: latent state variables evolve according to a Markov chain/Markov process
- Linear dynamical system: latent state variables evolve according to linear dynamics, possibly with added noise

What these have in common: hidden/latent state variables

# Hidden state models



- Latent/hidden system state
- Observations based on system state

# Two state-observation models

Two common models:

- Hidden Markov model

Hidden states evolve according to a Markov process

Observations typically Gaussian or multinomial

- Linear Gaussian dynamical system

States evolve according to linear dynamics

Observations a linear function of the state, “corrupted” by Gaussian noise

## Typical inference problems

Typical problems we want to solve, given a sequence of observations  $\mathcal{O}$  of time length  $T$ :

- Filtering: find the distribution of  $X_T$  – that is, the distribution of the current state, accounting for all observations up to now.
- Prediction: find the distribution of  $X_t$  for some  $t > T$ .
- Smoothing: find the distribution of  $X_t$  for some  $1 \leq t < T$ .  
This looks very similar to filtering, but differs in that we can take the observations after time  $t$  into account.
- MAP or best-explanation: find the sequence  $(X_i)$  maximizing  $P(\mathcal{O}, X)$ .
- Fitting: Given a sequence of observations, estimate the parameters of the underlying dynamical model.

# Hidden Markov models

---



## Example: the unfair casino

A casino employee has two 6-sided dice. We'll assume we know their properties:

Die	$P(1)$	$P(2)$	$P(3)$	$P(4)$	$P(5)$	$P(6)$
fair	$1/6$	$1/6$	$1/6$	$1/6$	$1/6$	$1/6$
loaded	$1/2$	$1/10$	$1/10$	$1/10$	$1/10$	$1/10$

The operator throws a die, but you don't know which one. What is the probability the die is loaded, assuming it lands on 1? What if instead it lands on 3?

# Tracking fairness over time

Let's say we know a little more about this casino employee's habits.

- The employee always starts the game with the fair die
- Every so often, they secretly switch the die
- Note: this is not an independent choice of die per throw

Result of this: streaks of fair/loading die over time.

If we observe the result of the die rolls, can we infer when each die was in use?

# Hidden Markov models

A hidden Markov model deals with two sequences:

- a sequence of *states*: the un-observed variable, changing over time according to a Markov chain model
- a sequence of *observations*, or *emissions*: the observed variable, with a distribution based on the current state

In our example:

- the state is which die is currently in use
- the emission is the roll of the die

## Simplest case

In a HMM, the underlying states are governed by a Markov process.

Our simple example is a finite state, multinomial HMM:

- Underlying state  $X_t$  follows a Markov chain with  $N$  states
- Observed values  $\mathcal{O}_t$  follow a multinomial distribution conditional on  $X_t$

So the model is described by two matrices,  $A$  (transition matrix), and  $B$  (observation matrix).

To do calculations, we also need to assume a certain distribution  $\pi$  on the initial state  $X_1$ . As a shorthand, I'll use the notation  $\lambda = (A, B, \pi)$  to represent a choice of these parameters.

Reference: Stamp, *A Revealing Introduction to Hidden Markov Models*

## Simplest case

Assumptions:

- Initial state is always fair, so

$$\pi = (1, 0)$$

- Transition matrix:

$$A = \begin{pmatrix} 0.95 & 0.05 \\ 0.05 & 0.95 \end{pmatrix}$$

- Observation matrix:

$$B = \begin{pmatrix} 1/6 & 1/6 & 1/6 & 1/6 & 1/6 & 1/6 \\ 1/2 & 1/10 & 1/10 & 1/10 & 1/10 & 1/10 \end{pmatrix}$$

# Three algorithms

Today: standard algorithms for filtering, smoothing, and fitting:

1. Given a multinomial HMM  $\lambda$  and a sequence of observations  $\mathcal{O}$ , compute the distribution  $P(X_t|\lambda, \mathcal{O})$ .
2. Given a multinomial HMM  $\lambda$  and a sequence of observations  $\mathcal{O}$ , compute the probability distribution of  $X_t$  for some  $1 \leq t \leq T$ .
3. Given a sequence of observations  $\mathcal{O}$ , what is the multinomial HMM  $\lambda$  that maximizes the marginal likelihood  $P(\mathcal{O}|\lambda)$ ?

## Naïve filtering

It is clear that we can compute the joint probability of a particular sequence of states with the observed sequence:

$$P(X, \mathcal{O} | \lambda) = \pi_{X_1} \prod_{t=1}^T A_{X_{t-1}, X_t} B_{X_t, \mathcal{O}_t}$$

So, naïvely, we could compute this for all sequences of states, and then

$$P(X_t = x_i) = \sum_{\text{sequences with } X_t = x_i} P(X, \mathcal{O} | \lambda)$$

What's the problem?

## Naïve filtering

It is clear that we can compute the joint probability of a particular sequence of states with the observed sequence:

$$P(X, \mathcal{O} | \lambda) = \pi_{X_1} \prod_{t=1}^T A_{X_{t-1}, X_t} B_{X_t, \mathcal{O}_t}$$

So, naïvely, we could compute this for all sequences of states, and then


$$P(X_t = x_i) = \sum_{\text{sequences with } X_t = x_i} P(X, \mathcal{O} | \lambda)$$

What's the problem?  $N^T$  sequences – computationally infeasible for all but short sequences.



# The forward algorithm


This problem can be solved by the *forward algorithm*, which exploits the Markov property to marginalize recursively on the fly:



Let  $\alpha_t(x_i) = P(X_t = x_i, \mathcal{O} | \lambda)$

# The forward algorithm

This problem can be solved by the *forward algorithm*, which exploits the Markov property to marginalize recursively on the fly:




Let  $\alpha_t(x_i) = P(X_t = x_i, \mathcal{O} | \lambda)$

$$\alpha_t(x_i) = B_{x_i, \mathcal{O}_i} \sum_{j=1}^N \alpha_{t-1}(x_j) A_{x_j, x_i}$$

# Filtering result

To test, generate a sequence of states and observations, and run the forward algorithm:



## The backward pass

The smoothing problem asks us to calculate  $P(X_t = x_i, \mathcal{O})$  for some  $t < T$ . We could just solve the filtering problem by running the forward algorithm up to time  $t$ , but we would lose the information from future states.

Solution: do a backward pass too.

Let  $\beta_t(x_i) = P(\mathcal{O}_{t:T} | X_t = x_i)$ ; that is, the probability of the "remaining" observations from time  $t$  to the end, given  $X_t = x_i$ .

Then,

$$\beta_t(x_i) = \sum_{j=1}^N A_{x_i, x_j} B_{x_i, \mathcal{O}_t} \beta_{t+1}(x_j)$$

so we can recursively calculate from the end of the sequence, letting  $\beta_T(x_j) = 1$  for each  $j$ .

# The forward-backward algorithm

The forward-backward algorithm solves the smoothing problem for HMMs:


$$P(X_t = x_i | \mathcal{O}, \lambda) = \frac{\alpha_t(x_i)\beta_t(x_i)}{P(\mathcal{O}|\lambda)}$$

Where can we get the normalizing constant?

$$P(\mathcal{O}|\lambda) = \sum_{i=1}^N \alpha_T(x_i)$$

# Smoothing result

To test, generate a sequence of states and observations, and run the forward-backward algorithm:



# Fitting parameters

The fitting problem gives a new challenge:

- Given a fixed state space  $\{0, 1, \dots, n\}$  and a sequence  $\mathcal{O}$  of observations, find the model parameters that best fit the sequence  $\mathcal{O}$
- i.e., tune  $A$  (transition matrix),  $B$  (observation matrix), and  $\pi$  (initial state distribution)
- Target: maximize  $P(\mathcal{O}|A, B, \pi)$

This is a form of unsupervised learning.

# Baum-Welch algorithm

The Baum-Welch algorithm iteratively improves the fit of the model parameters in a two-step process:

- Do a smoothing step, estimating the probability distributions of the hidden states  $X_t$
- Re-adjust the model parameters to better fit this estimated distribution
- Score the model by the log-probability of the observed sequence
- Continue until log-probability change is negligible

End result: MAP estimate of model parameters



# Idea behind BW algorithm

Intuitively:

- The smoothing step allows us to estimate the probability that the underlying chain is in each state  $x_i$  at time  $t$
- We can use this to count the estimated probability of transitions from state  $x_i$  to state  $x_j$
- We can use this, together with the observation sequence, to estimate the probability of each observation from state  $x_i$

## Estimating the transition matrix

Smoothing gives us:

$$\gamma_t(i) = \frac{\alpha_t(i)\beta_t(i)}{P(\mathcal{O}|A, B, \pi)}$$

which estimate the probability that the chain was in state  $x_i$  at time  $t$ . We extend this to:

$$\gamma_t(i, j) = \frac{\alpha_t(i)A_{ij}B_{j, \mathcal{O}_t} \beta_{t+1}(j)}{P(\mathcal{O}|A, B, \pi)}$$

which estimates the probability that the chain was in state  $x_i$  at time  $t$  *and* state  $x_j$  at time  $t + 1$ .

Then, we re-estimate the transition probability  $A_{ij}$  as:

$$A_{ij} = \frac{\sum_t \gamma_t(i, j)}{\sum_t \gamma_t(i)}$$

Similarly, we can re-estimate the observation probability  $B_{ij}$  as

$$B_{ij} = \frac{\sum_{t, \mathcal{O}_t=j} \gamma_t(i)}{\sum_t \gamma_t(i)}$$

the expected proportion of the time spent in state  $i$  that produces observation  $j$ .

## Estimating the rest

Similarly, we can re-estimate the observation probability  $B_{ij}$  as

$$B_{ij} = \frac{\sum_{t, \mathcal{O}_t=j} \gamma_t(i)}{\sum_t \gamma_t(i)}$$

the expected proportion of the time spent in state  $i$  that produces observation  $j$ .

The estimate of the initial state vector is just:

$$\pi_i = \gamma_0(i)$$


# Testing the algorithm

Let's test the algorithm on the unfair casino problem:

- Generate 1200 observations from the “true” model
- Initialize a HMM with the correct number of states, but randomly initialized parameters
- Fit the model; test its performance on a smoothing problem


# Testing the algorithm

Now: 200 new states and observations



# Testing the algorithm

Another run of this experiment produced this:



Is this a failure of the fitting process?

## Aside: Using fitted HMM for recognition problems

Old application: speech recognition

- Given a set of phrases you want to recognize:
  - Fit a different HMM to example data of each phrase
  - To recognize a new speech segment, compute the log-probability  $\log P(\mathcal{O})$  with each HMM
- Fitting process maximizes  $\log P(\mathcal{O})$ , so a model trained on a specific phrase should assign a high log probability to that phrase
- These days, supplanted by recurrent neural networks



## Aside 2: Viterbi algorithm

The forward-backward algorithm computes

$$P(X_t = i | \mathcal{O}, \lambda)$$

for each  $1 \leq t \leq T$ , so we can get a MAP estimate of the state at each fixed time.

- Slightly different: MAP estimate of the state sequence
- Most probable sequence  $X$  may not equal the most probable state at each time step
- Viterbi algorithm (described in Stamp's article as *dynamic programming*) gives MAP estimate for the state sequence

# Viterbi algorithm

Like the forward algorithm, the Viterbi algorithm exploits the Markov property to express the calculation recursively:

- Assume we can calculate the most probable sequences that end in states  $x_i$  at time  $t - 1$
- The most probable sequence that ends in state  $x_j$  at time  $t$  must be an extension of one of the  $N$  most probable sequences up to time  $t - 1$
- The most probable sequence overall must be the most probable sequence ending in state  $x_i$  at time  $T$ , for some  $i$

# Viterbi algorithm

Essentially, we modify the forward algorithm to replace a sum with a max. Let:

$$\tilde{\alpha}_t(x_i) = P(X_{1:t} = \tilde{X}_{1:t,i}, \mathcal{O}_{1:t} | \lambda)$$

where  $\tilde{X}_{1:t-1,i}$  is the most probable sequence of steps given  $\mathcal{O}$  such that  $X_t = x_i$ .

$$\tilde{\alpha}_t(x_i) = \max_j (A_{x_j, x_i} B_{x_i, \mathcal{O}_t} \tilde{\alpha}_{t-1}(x_j))$$

## Other applications

Because of their ability to find and recognize patterns in sequence data without supervision, HMMs find applications in, among other places:

- speech recognition
- cryptanalysis (codebreaking) and malware detection
- activity recognition
  - identify what an animal is doing based on GPS data

## Text analysis example

Imagine you're an alien with no knowledge of human language, but you gain access to a sample of English text, and you would like to extract some information about the relationships between characters.

Simplifying assumptions:

- No cases – everything is lowercase
- No digits or punctuation; only characters are letters and spaces

Idea:

- different characters play different roles in the written language.
- fit a hidden Markov model with  $k$  different states to a large sample of text, and see if any patterns can be seen.

Let's take a look at the results for  $k = 2$ .

# Expectation-maximization algorithms

---

# EM algorithms

The Baum-Welch algorithm we saw before is an example of a much wider class of algorithms called *expectation-maximization* algorithms.

These are applicable when the observed data depends on hidden/latent state variables as well as model parameters. Roughly, the idea is:

- Expectation step: compute the distribution of hidden state variables, given current model parameters
- Maximization step: compute the model parameters that maximize (log) likelihood given the state parameters from the expectation step

Repeat until done – score model by total log-likelihood of the data.

Section 13.4-13.6 in BDA has another presentation of EM algorithms in a different context.

Formally:

- $\theta$ : model parameters
  - $X$ : hidden variables
  - $Y$ : observations
  - $L(\theta|X, Y)$ : likelihood function
1. E-step: compute  $Q(\theta|\hat{\theta}) = E_{X|Y, \hat{\theta}}[\log L(\theta|X, Y)]$
  2. M-step: compute  $\theta^{\text{new}} = \arg \max_{\theta} Q(\theta|\hat{\theta})$



Recall the Baum-Welch algorithm has two steps:

- Perform smoothing to estimate the distribution of each  $X_t$ , given current transition/observation matrix values
- Update parameter values by counting transitions/observations given distributions of  $X_t$

Although we don't explicitly calculate expectations of log-likelihoods, the smoothing step is an E step and the update step is an M step.

## **A few comments**

---

# Alternative approaches to parameter estimation

More Bayesian approaches to estimating HMM parameters:

- MCMC – suitable for smaller systems
  - Helmut Strey has a PyMC3 implementation
  - <https://github.com/hstrey/Hidden-Markov-Models-pymc3>
  - I haven't checked this out carefully
- Particle filter (aka sequential Monte Carlo) methods
  - Can also replace Kalman filters for nonlinear dynamical models
  - We'll return to this

The most common alternative to the multinomial distribution for HMMs is the Gaussian (normal) distribution. In this model:

- $X_t$  still evolves according to a Markov chain with transition matrix  $A$
- $\mathcal{O}_t \sim \text{MVNormal}(\mu_{X_t}, \Sigma_{X_t})$
- Result: the observation distributions are Gaussian mixtures

What has to change for our filtering and smoothing algorithms?

What has to change for our filtering and smoothing algorithms?

- Only change:  $P(\mathcal{O}_t = j | X_t = x_i)$  is no longer given by a matrix entry  $B_{ij}$
- Instead, we have  $p(\mathcal{O}_t = y | X_t = x_i) = \text{MVNormal}(\mu_i, \Sigma_i)$  for a certain mean vector  $\mu_i$ , covariance matrix  $\Sigma$

# EM for Gaussian HMM

To fit the Gaussian HMM, we only need to make the following modifications to the M step:

- Replace  $B_{ij}$  with  $\mu_i, \Sigma_i$
- Replace the update of  $B_{ij}$  with a maximum-likelihood estimate for a Gaussian, weighted by the estimated state probabilities (from smoothing):

$$\mu_i^{\text{new}} = \frac{\sum_t P(X_t = i) \mathbf{y}_t}{\sum_t P(X_t = i)}$$
$$\Sigma_i^{\text{new}} = \frac{\sum_t P(X_t = i) (\mathbf{y}_t - \mu_i^{\text{new}})(\mathbf{y}_t - \mu_i^{\text{new}})^T}{\sum_t P(X_t = i)}$$

where  $\mathbf{y}_t$  are the observations.

Today:

- Hidden Markov models

Next week: Rethinking chapters 13, 15

- Mixture models
- Missing data